

# Feature Engineering with SKLearn-Pandas

---

By Ramesh Sampath (@sampathweb)

Slides: [bit.ly/sk-pandas](https://bit.ly/sk-pandas)

# Ramesh Sampath

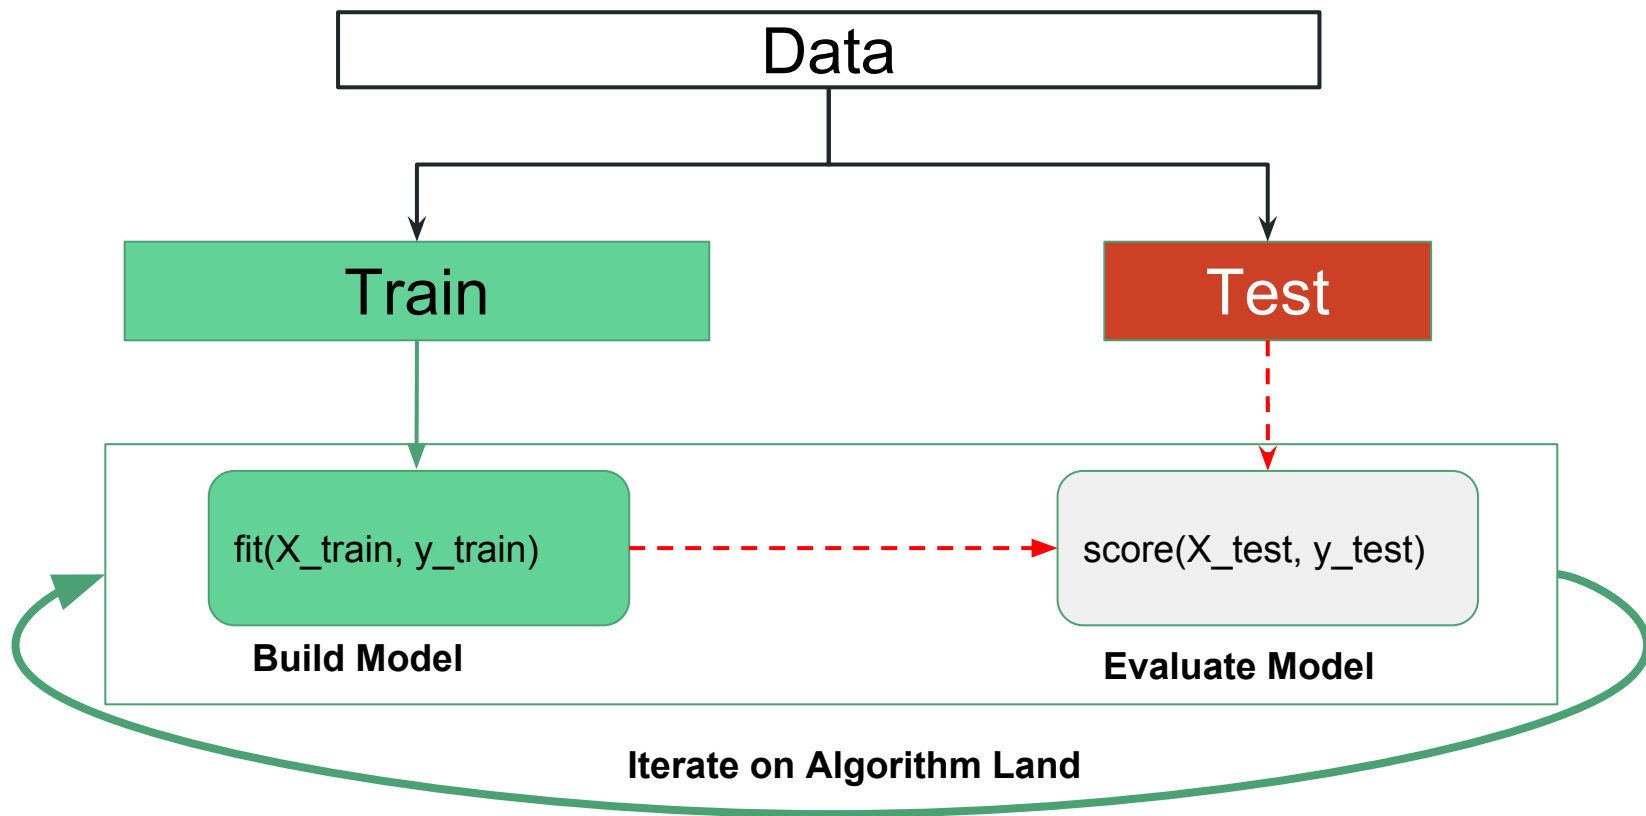
- Data Science Engineer
  - Some Machine Learning Models
  - A lot of Pre-Processing
  - Deploy it as API Services

@sampathweb (github / twitter / linkedin)

# What's the Problem

- Data Scientists Want to -
  - Build Models
  - Tune Models
  - Spend time in Algorithm Land

# Ideal World



# Data is Messy

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25	NaN	S
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.28	C85	C

Data needs to be Numerical Vector for Machine Learning.

$$\frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

# Vectorizing

Survived	Pclass	Sex	Age	SibSp	Embarked
0	3	male	22.0	1	S
1	1	female	38.0	1	C

**Target -  
Classification**

**Class -  
Categorical**

**Gender -  
Categorical**

**Age -  
Continuous, N/A**

**Sibling -  
Count**


**Embarked -  
Categorical, N/A**

Age	SibSp	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
22.0	1	0	0	1	0	1	0	0	1
38.0	1	1	0	0	1	0	1	0	0
26.0	0	0	0	1	1	0	0	0	1

$$\frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Logistic Regression

Data (Numerical / Categorical / Missing Values)



SKLearn-Pandas

Clean, Numerical Data Matrix



`fit(X_train, y_train)`

Scikit-Learn

# SKLearn-Pandas



- Original code by Ben Hamner (Kaggle CTO) and
- Paul Butler (Google NY) 2013

```
from sklearn_pandas import DataFrameMapper
```



# SKLearn-Pandas

```
mapper = DataFrameMapper([
    (["Pclass"], OneHotEncoder()),
    (["Sex"], LabelBinarizer()),
    (["Age", "Fare"], [Imputer(strategy="mean"), StandardScaler()]),
    ("Embarked", [ColumnImputer(strategy="most_frequent"), LabelBinarizer()]),
    (["SibSp", "Parch"], None) # No Transformations necessary
])

pipeline = Pipeline(steps=[
    ("features", mapper),
    ("model", LogisticRegression())
])

pipeline.fit(X_train, y_train)

print("Accuracy Score on Test Data: {:.3f}".format(pipeline.score(X_test, y_test)))
```

Accuracy Score on Test Data: 0.825

# Feature Engineering Pipeline

## Pre-Processing

- Cleaning / Imputing Values
- Encoding to Numerical Vectors

## Feature Extractions

- Text Vectorization (Count / TFIDF)
- Polynomial Features

## Feature Reduction & Selection

- PCA
- SelectFromModel

Machine Learning Models

**Grid Search - Hyper Parameter Tuning of Models**

# Grid Search

```
pipeline = Pipeline(steps=[
    ("features", mapper),
    ("rf_model", RandomForestClassifier())
])
```

```
params = {
    "rf_model__n_estimators": [10, 100, 200],
    "rf_model__min_samples_leaf": [1, 3, 5],
    "rf_model__max_depth": [None, 10, 7, 5, 3]
}
```

Hyper Parameter Tuning (Hurray!)  
Back in Algorithm Land

```
grid = RandomizedSearchCV(pipeline, param_distributions=params, n_iter=5, n_jobs=-1)

grid.fit(X_train, y_train)
print("Best Parameters: ", grid.best_params_)
print("Accuracy Score on Test Data: {:.3f}".format(grid.score(X_test, y_test)))
```

# Data is Messy

age	workclass	education	education_num	marital_status	occupation	relationship	race	sex	native_country	income
30	Private	Some-college	10	Never-married	Sales	Own-child	White	Male	United-States	<=50K
34	Private	HS-grad	9	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	<=50K
41	Private	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	United-States	>50K
58	Private	10th	6	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States	<=50K

# Pipeline

```
mapper = DataFrameMapper([
    ("workclass", LabelBinarizer()),
    ("marital_status", LabelBinarizer()),
    ("occupation", LabelBinarizer()),
    ("relationship", LabelBinarizer()),
    ("race", LabelBinarizer()),
    ("sex", LabelBinarizer()),
    ("native_country", LabelBinarizer()),
    ("workclass", LabelBinarizer()),
    ("workclass", LabelBinarizer()),
    ([ "age", "education_num"], None) # No Transformations necessary
])

pipeline = Pipeline(steps=[
    ("features", mapper),
    ("model", LogisticRegression())
])

pipeline.fit(X_train, y_train)

print("Accuracy Score on Test Data: {:.3f}".format(pipeline.score(X_test, y_test)))
```

Accuracy Score on Test Data: 0.831

# Jupyter Notebook

<https://github.com/sampathweb/odsc-feature-engineering-talk>

# Credits

- Scikit-Learn (<https://github.com/scikit-learn/scikit-learn>)
- Sklearn-Pandas (<https://github.com/paulgb/sklearn-pandas>)

## StackOverflow Posts:

- <http://stackoverflow.com/questions/24458645/label-encoding-across-multiple-columns-in-scikit-learn>
- <http://stackoverflow.com/questions/34710281/use-featureunion-in-scikit-learn-to-combine-two-pandas-columns-for-tfidf>

# Thank You!

**Slides:** [bit.ly/sk-pandas](https://bit.ly/sk-pandas)

@sampathweb (Github / Twitter / LinkedIn)